

# CSCI 4022 Fall 2021 Community Detection



The **AGM model**: generates the links of a network *probabilistically*, using the community probabilities  $p_C$ .

1. For each pair of nodes  $u, v$  in a community  $A \in C$ , we connect them with probability  $p_A$ .
2. ... but  $u$  and  $v$  might also share e.g. community  $B \in C$ , which would *independently* connect them with probability  $p_B$ .  $P(A-B)$ : edge  $\leftarrow$  from  $0$  or  $1$   $p$   $q$
3. In the final network  $u$  and  $v$  will be connected by an edge if an edge is generated from *any one* of their shared communities

What's the overall probability that *at least one* such edge is drawn?

Instead:  $P(\text{no } A-B) = P(\text{not from any community they share})$

		$p$	$q$	$P$
I	T	$p \cdot q$		
	F	$p(1-q)$		
T	T	$(1-p)q$		
	F	$(1-p)(1-q)$		

## Announcement and Reminders

+ Z office hour 5p-6p

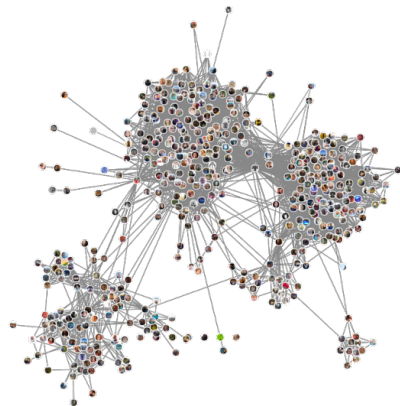
1. Hw 5 due tonight HW6 **and** project proposals for next Monday.

# Networks and Communities

## Examples:

Facebook friends.

1. Nodes: users
2. Edges: friendships



# Networks and Communities

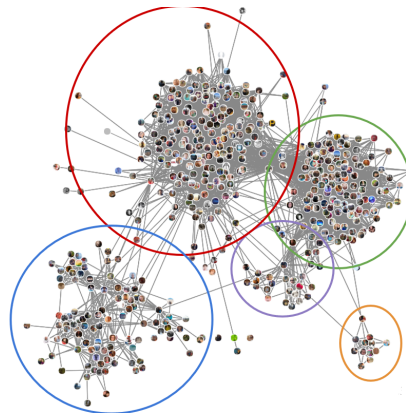
## Examples:

Facebook friends.

1. Nodes: users
2. Edges: friendships

This also invites the notion of a community:

1. red: College friends
2. green: high school friends
3. blue: graduate school friends
4. purple: family
5. orange: summer internship friends

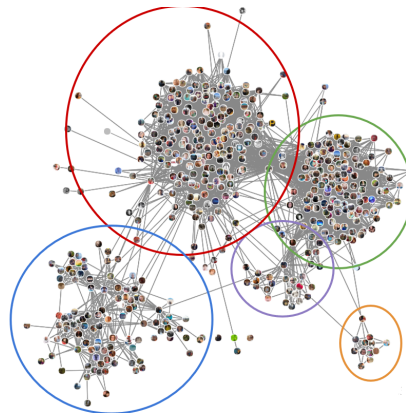


# Networks and Communities

## Examples:

Facebook friends.

1. Nodes: users
2. Edges: friendships



What could we do with this?

Describe points with *similar* community affiliations: recommend friends, events, or advertisements.

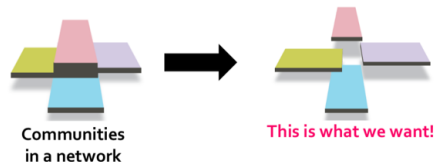
Unlike clusters, these groups **overlap**, and nodes will belong to multiple groups.

# Communities

**Goal:** Find a way to describe the communities in a graph.

**Similarity** to clustering:

1. There is overlap between different “clusters” – each is a community
2. Unlike clustering, can belong to multiple communities
3. The goal: can we identify these social network communities?



## Modeling a Network and Communities

If we can come up with a **model**, we could make a graph. Think of this as the conditional problem: *given* communities, how might we draw a graph?

**Goal:** given a model, generate networks

1. Given some nodes. . .
2. The model will have a set of parameters that govern the connections among nodes
3. We will estimate these parameters (today, next time). But if we have a generative model that's based on *probability*, than the “best choice of parameters,” sounds in all *likelihood* to be a problem we've dealt with before...

Question: Given a set of nodes, how will our model generate the edges of the network?

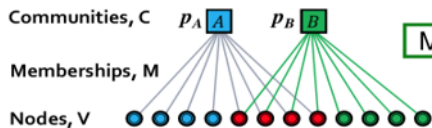


# The Community-Affiliation Graph Model (AGM)

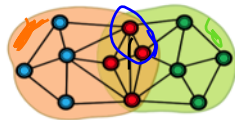
**Model:** Given a set of nodes, what's a reasonable way to generate edges of a network?

The **AGM model:**  $B(V, C, M, \{p_c\})$

1.  $V$  = the set of nodes.
2.  $C$  = the set of communities.
3.  $M$  = the set of memberships (think of it as edges from  $V$  to  $C$ , or object that's  $\text{len}(V) \times \text{len}(C)$  and holds a score if row "person" is inside column "community".)
4. Each community  $C$  gets a unique probability  $p_C$  denoting the probability of  $C$  generating a connection between members of  $C$ .



Model



Network

		Comm.		
		1	2	3
A	0	1	2	3
	1	1	1	0
B	1	0	1	1

$$P(A-B \text{ exists}) = 1 - (1-p_0)(1-p_2)$$

not from 0

$$P(\text{edge}) = 1 - (1-p)(1-q)$$

$$P(\text{not edge}) = (1-p) \cdot (1-q)$$

term for each community shared



## The Community-Affiliation Graph Model (AGM)

The **AGM model**: generates the links of the network *probabilistically*, using the community probabilities  $p_C$ .

1. For each pair of nodes  $u, v$  in a community  $A \in C$ , we connect them with probability  $p_A$ .
2. ... but  $u$  and  $v$  might also share e.g. community  $B \in C$ , which would *independently* connect them with probability  $p_B$ .
3. In the final network  $u$  and  $v$  will be connected by an edge if an edge is generated from *any one* of their shared communities  $\rightarrow$  (one or more) = exactly 0.

What's the overall probability that at least one such edge is drawn? **DeMorgan's Laws**

$P(\text{at least one edge}) = 1 - P(\text{no edges}) \dots$

$= 1 - P(\text{no edge from first shared comm AND no edge from first shared comm AND} \dots)$

**Result:** under AGM, the probability of the edge  $(u, v)$  is

$$P(u, v) = 1 - \prod (1 - p_c).$$

$C \in M_u \cap M_v$   $\rightarrow$  shared communities

## The Community-Affiliation Graph Model (AGM)

The **AGM model**: generates the links of the network *probabilistically*, using the community probabilities  $p_C$ .

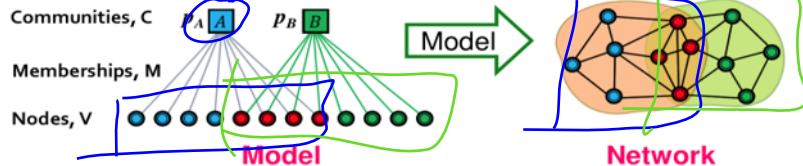
**Result:** under AGM, the probability of the edge  $(u, v)$  is

$$P(u, v) = 1 - \prod_{C \in M_u \cap M_v} (1 - p_C).$$

where  $M_u$  is the set of all communities of member  $u$ .

We may also include a *background probability*, where

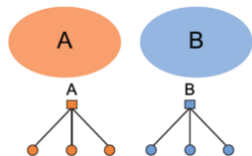
$P(u, v) = \epsilon$  if  $M_u \cap M_v = \emptyset$  for some small  $\epsilon$ . What does this  $\epsilon$  represent?



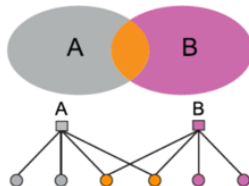
# Flexibility of AGM

We can use AGM to express a variety of different common structures:

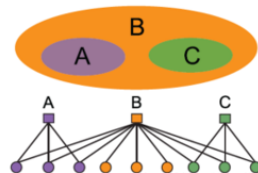
**Non-overlapping; disjoint**



**Overlapping**

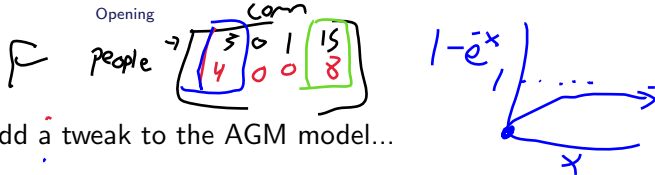


**Nested (Subset) Communities**



Now, to actually estimate those probabilities *given a graph*...

## Modifying AGM



For two reasons, we're going to add a tweak to the AGM model...

1. Membership in communities is not binary (0/1), it has a strength. (The team captain is more likely to form within-team connections than a supporting member)

Define:  $F_{u,A}$ : the membership strength of node  $u$  in community  $A$ . Set  $F_{u,a} = 0$  if and only if  $u$  is absolutely **not** a member of  $A$ .

2. We adjust the resulting probability to link two nodes in the same community  $A$  to be

for each community

$$P_A(u, v) = 1 - \exp(-F_{u,A} \times F_{v,A}).$$

$1 - e^{-(0.3 \cdot 4)}$

Why this? If  $u$  and  $v$  share many groups, the probability of the edge  $(u, v)$  given by

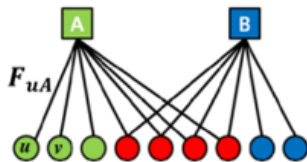
$$P(u, v) = 1 - \prod_{C \in M_u \cap M_v} (1 - p_C).$$

not from any

is going to be **much easier** if it simplifies!

# BigCLAM

The resulting model is called the Cluster Affiliation Model for Big Networks, or BigCLAM.



$F_{u,A}$ : the membership strength of node  $u$  in community  $a$ . Set  $F_{u,a} = 0$  if and only if  $u$  is absolutely **not** a member of  $A$ .

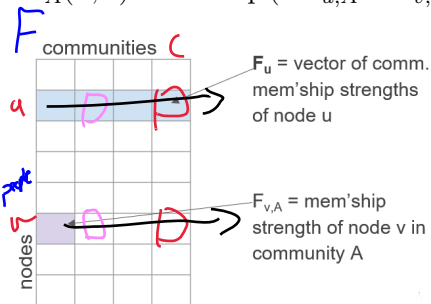
$$P_A(u, v) = 1 - \exp(-F_{u,A} \times F_{v,A}).$$

**Goal:** Find the  $k$  communities inside a *given* graph... kind of like a GMM or k-means!

# BigCLAM

So we've replaced the  $p_c$ 's of AGM with a new probability-per-community that relies on membership strengths:

$$P_A(u, v) = 1 - \exp(-F_{u,A} \times F_{v,A}).$$



Prob (graph |  $F$ )

...This helps with the probability that nodes  $u$  and  $v$  share *at least one edge* through any community:

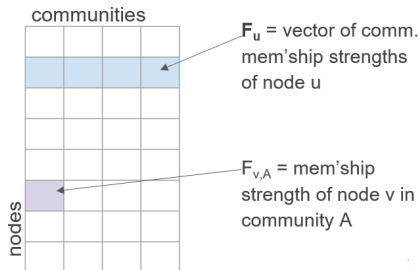
$$P(u, v) = 1 - \prod_c (1 - P_c(u, v))$$

$\underbrace{1 - e^{-(F_{u,c} \cdot F_{v,c})}}$

# BigCLAM

So we've replaced the  $p_c$ 's of AGM with a new probability-per-community that relies on membership strengths:

$$P_A(u, v) = 1 - \exp(-F_{u,A} \times F_{v,A}).$$



$$e^a \cdot e^b \cdot e^c = e^{a+b+c}$$

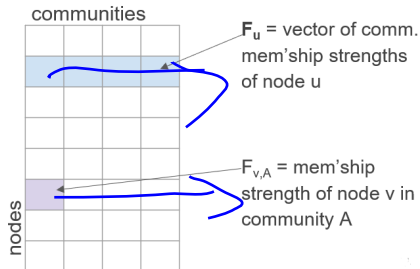
...This helps with the probability that nodes  $u$  and  $v$  share *at least one edge* through any community:

$$\begin{aligned} P(u, v) &= 1 - \prod_c (1 - P_c(u, v)) \\ &= 1 - \prod_c (1 - (1 - \exp(-F_{u,c} \times F_{v,c}))) \\ &= 1 - \prod_c \exp(-F_{u,c} \times F_{v,c}) \end{aligned}$$

# BigCLAM

So we've replaced the  $p_c$ 's of AGM with a new probability-per-community that relies on membership strengths:

$$P_A(u, v) = 1 - \exp(-F_{u,A} \times F_{v,A}).$$



...This helps with the probability that nodes  $u$  and  $v$  share *at least one edge* through any community:

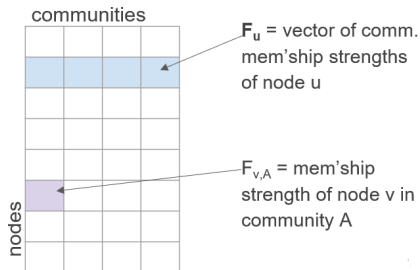
$$\begin{aligned} P(u, v) &= 1 - \prod_c (1 - P_c(u, v)) \\ &= 1 - \prod_c (1 - (1 - \exp(-F_{u,c} \times F_{v,c}))) \\ &= 1 - \prod_c \exp(-F_{u,c} \times F_{v,c}) \\ &= 1 - \exp\left(\sum_c -F_{u,c} \times F_{v,c}\right) \end{aligned}$$



# BigCLAM

So we've replaced the  $p_c$ 's of AGM with a new probability-per-community that relies on membership strengths:

$$P_A(u, v) = 1 - \exp(-F_{u,A} \times F_{v,A}).$$



...This helps with the probability that nodes  $u$  and  $v$  share *at least one edge* through any community:

$$\begin{aligned}
 P(u, v) &= 1 - \prod_c (1 - P_c(u, v)) \\
 &= 1 - \prod_c (1 - (1 - \exp(-F_{u,c} \times F_{v,c}))) \\
 &= 1 - \prod_c \exp(-F_{u,c} \times F_{v,c}) \\
 &= 1 - \exp\left(\sum_c -F_{u,c} \times F_{v,c}\right) \\
 &= 1 - \exp(-F_u \cdot F_v)
 \end{aligned}$$

*Handwritten notes: "not edge" above the first product, "not from c" next to the second product.*

## BigCLAM Probabilities

$$P(\text{edge}_{u,v}) = 1 - e^{-(F_u \cdot F_v)}$$

Example: Suppose for node set  $V = \{u, v, w\}$  and communities  $\{A, B, C, D\}$ , we have the following membership strength matrix  $F$ . What are the probabilities of a connection between  $u$  and  $v$ ? Between  $u$  and  $w$ ?  $v$  and  $w$ ?

$$F = \begin{array}{c|cccc} & A & B & C & D \\ \hline u & 0 & 1.2 & 0 & 0.2 \\ v & 0.5 & 0 & 0 & 0.8 \\ w & 0 & 1.8 & 1 & 0 \end{array}$$

$$P(\text{edge}_{u,v}) = 1 - e^{-(F_u \cdot F_v)}$$

$$P(u,v) = 1 - e^{-(F_u \cdot F_v)} = 1 - e^{-(0.2 \cdot 0.8)} = 1 - e^{-0.16}$$

$$P(u,w) = 1 - e^{-(0 \cdot 0 + 1.2 \cdot 1.8 + 0 \cdot 1 + 0.2 \cdot 0)} = 1 - e^{-(1.2 \cdot 1.8)}$$

$$P(v,w) = 1 - e^{-0} = 1 - 1 = 0$$

# BigCLAM Probabilities

Example: Suppose for node set  $V = \{u, v, w\}$  and communities  $\{A, B, C, D\}$ , we have the following membership strength matrix  $F$ . What are the probabilities of a connection between  $u$  and  $v$ ? Between  $u$  and  $w$ ?  $v$  and  $w$ ?

$$F =$$

	A	B	C	D
$u$	0	1.2	0	0.2
$v$	0.5	0	0	0.8
$w$	0	1.8	1	0

$$P(\text{edge}_{u,v}) = 1 - e^{-(F_u \cdot F_v)}$$

**Solution:**

$$P(u, v) = 1 - e^{-(F_u \cdot F_v)} = 1 - e^{-0.16} = 0.14$$

$$P(u, w) = 1 - e^{-(F_u \cdot F_w)} = 1 - e^{-2.16} = 0.88$$

$$P(v, w) = 1 - e^{-(F_v \cdot F_w)} = 0$$

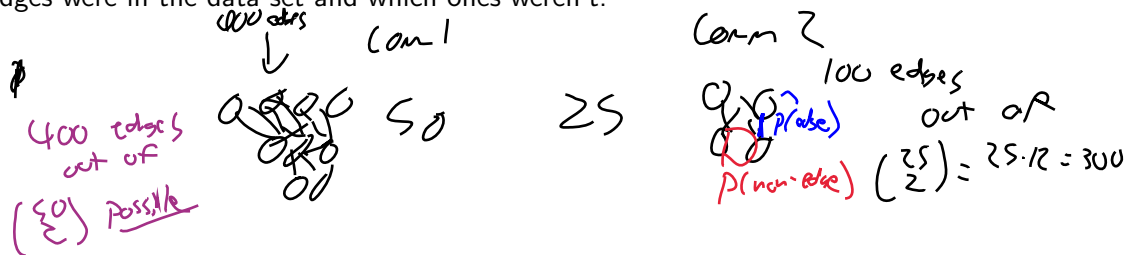
## BigCLAM Implementation

So connection probabilities are pretty easy, as  $P(u, v) = 1 - \exp(-F_u \cdot F_v)$  ...As long as we know  $F$ . Given the undirected graph  $G(V, E)$ , how do we find  $F$ ?

## BigCLAM Implementation

So connection probabilities are pretty easy, as  $P(u, v) = 1 - \exp(-F_u \cdot F_v)$  ... As long as we know  $F$ . Given the undirected graph  $G(V, E)$ , how do we find  $F$ ?

**Answer:** Pick  $F$  to maximize the *likelihood function*. Recall that a likelihood function is the probability of the data *given* the model, or the probability that we observed exactly which edges were in the data set and which ones weren't.



## BigCLAM Implementation

So connection probabilities are pretty easy, as  $P(u, v) = 1 - \exp(-F_u \cdot F_v)$  ...As long as we know  $F$ . Given the undirected graph  $G(V, E)$ , how do we find  $F$ ?

**Answer:** Pick  $F$  to maximize the *likelihood function*. Recall that a likelihood function is the probability of the data *given* the model, or the probability that we observed exactly which edges were in the data set and which ones weren't.

The edge connections between nodes are a Bernoulli process:

- ▶  $(u, v) \in E$  with probability  $P(u, v)$  or  $(u, v) \notin E$  with probability  $1 - P(u, v)$
- ▶ Each edge connection is established probabilistically, independently of the others.

**Result:** Probability of a given graph (edge set  $E$ ) being established, given  $F$ , is:

$$L(F) = \prod_{(u,v) \in E} P(u, v) \prod_{(u,v) \notin E} (1 - P(u, v)) \quad \leftarrow \binom{2}{2} \text{ things}$$

## BigCLAM Likelihood

**Result:** Probability of a given graph (edge set  $E$ ) being established, given  $F$ , is:

$$L(F) = \prod_{(u,v) \in E} P(u,v) \prod_{(u,v) \notin E} (1 - P(u,v))$$

**Goal:** find the community affiliations  $F$  that maximize this value.

Here's the thing about likelihood functions: nobody likes them. Nobody. He's an unstable numerical mess.

Likelihood Function's little sister, the log-likelihood function, on the other hand, everybody likes:

$$l(F) = \log \left( \prod_{(u,v) \in E} P(u,v) \prod_{(u,v) \notin E} (1 - P(u,v)) \right)$$

$$l(F) = \sum_{(u,v) \in E} (\log [1 - \exp(-F_u \cdot F_v)]) - \sum_{(u,v) \notin E} (\exp(-F_u \cdot F_v))$$

# BigCLAM Log-Likelihood

**Goal:** find the community affiliations  $F$  that maximize the log-likelihood function:

$$l(F) = \log \left( \prod_{(u,v) \in E} P(u,v) \prod_{(u,v) \notin E} (1 - P(u,v)) \right)$$

$$l(F) = \sum_{(u,v) \in E} (\log [1 - \exp(-F_u \cdot F_v)]) + \sum_{(u,v) \notin E} \log (1 - 1 - \exp(-F_u \cdot F_v))$$

$$l(F) = \underbrace{\sum_{(u,v) \in E} (\log [1 - \exp(-F_u \cdot F_v)])}_{\text{edges}} - \underbrace{\sum_{(u,v) \notin E} F_u \cdot F_v}_{\text{no edges}}$$

now... **commence maximization!!**



## BigCLAM Log-Likelihood

**Goal:** estimate  $F$  by optimizing the log-likelihood function of

$$l(F) = \sum_{(u,v) \in E} \log(1 - \exp(-F_u \cdot F_v)) - \sum_{(u,v) \notin E} F_u \cdot F_v$$



via *gradient ascent*.

Denote  $N(u)$  by the set of neighbor nodes of  $u$ , which are connected to  $u$  by edges but may or may not share many communities. **Idea:**

1. The log-likelihood  $l(F)$  defines a surface with respect to the “coordinates” (rows of  $F$ )
2. We want to ascend to the top (maximum) of the log-likelihood surface
3. Fun fact: the gradient points uphill
4. Given a particular guess for a row in  $F$ ,  $F_u$ , we can improve it by taking a step in the direction of the gradient of  $l(F)$  with respect to (vector)  $F_u$ :  $\nabla l(F_u)$

## Gradient Ascent/Descent

(local)



The idea of following the derivative to find the maximum or minimum value of a function is called gradient ascent or gradient descent. It requires:

1. The ability to calculate the slope of the function we're trying to min/max
2. An idea of how large of steps to take; a step size or *learning rate*

$$\underbrace{F^{(k+1)}}_{\text{new guess}} = \underbrace{F^{(k)}}_{\text{old guess}} + \underbrace{\nu}_{\text{step size}} \underbrace{F'(z^{(k)})}_{\text{step direction}}$$

In our case, we're going to *update* our estimates for  $F$  by taking small steps in the direction of the community affiliations for node  $u$ . In other words: A step is an update to  $u$  to be more closely affiliated with it's neighbors. Then repeat for every node  $u$ .

The *gradient* is the multivariate direction we're supposed to take steps in!

## Gradient Ascent/Descent

So we need a derivative to determine which way to take a step.

**Idea:** move community affiliations of a node closer to the affiliations of its neighbors.

In practice, we're differentiating

$$l(F) = \sum_{(u,v) \in E} \log(1 - \exp(-F_u \cdot F_v)) - \sum_{(u,v) \notin E} F_u \cdot F_v$$

but we'll go at it *one specific location* at a time, so we're looking at

$$l(F_u) = \sum_{v \in N(u)} \log(1 - \exp(-F_u \cdot F_v)) - \sum_{v \notin N(u)} F_u \cdot F_v$$

and differentiating with respect to row  $u$

(In other words “how should we update our knowledge of person  $u$ ”).

**Calculus friends:**  $\frac{d}{dx} \log(1 - f(x)) =$  ,

$$\frac{d}{dx} e^{f(x)} =$$

## Gradient Ascent/Descent

So we need a derivative to determine which way to take a step.

**Idea:** move community affiliations of a node closer to the affiliations of its neighbors.

In practice, we're differentiating

$$l(F) = \sum_{(u,v) \in E} \log(1 - \exp(-F_u \cdot F_v)) - \sum_{(u,v) \notin E} F_u \cdot F_v$$

but we'll go at it *one specific location* at a time, so we're looking at

$$l(F_u) = \sum_{v \in N(u)} \log(1 - \exp(-F_u \cdot F_v)) - \sum_{v \notin N(u)} F_u \cdot F_v$$

and differentiating with respect to row  $u$

(In other words “how should we update our knowledge of person  $u$ ”).

**Calculus friends:**  $\frac{d}{dx} \log(1 - f(x)) = \frac{f'(x)}{1 - f(x)},$

$$\frac{d}{dx} e^{f(x)} = f'(x) e^{f(x)}$$

# The BigCLAM gradient

$$\nabla l(F_u) = \frac{d}{dF_u} \sum_{v \in N(u)} \log(1 - \exp(-F_u \cdot F_v)) - \sum_{v \notin N(u)} F_u \cdot F_v$$

Each term in the first sum is a derivative of  $\log(1 - \exp(-F_u \cdot F_v))$ , which gives  $F_v \frac{\exp(-F_u \cdot F_v)}{1 - \exp(-F_u \cdot F_v)}$ .

Each term in the second sum is a derivative of  $F_u \cdot F_v$ , so we are left with just  $F_v$ .

**Result:**

$$\nabla l(F_u) = \left\langle \underbrace{\sum_{v \in N(u)} F_{v,A} \frac{\exp(-F_u \cdot F_v)}{1 - \exp(-F_u \cdot F_v)} - \sum_{v \notin N(u)} F_{v,A}, \dots}_{\nabla_A l(F_u)} \right\rangle$$

# The BigCLAM gradient

## The full update:

$$\begin{aligned} \nabla l(F_u) = \langle & \sum_{v \in N(u)} F_{v,A} \frac{\exp(-F_u \cdot F_v)}{1 - \exp(-F_u \cdot F_v)} - \sum_{v \notin N(u)} F_{v,A}, \\ & \sum_{v \in N(u)} F_{v,B} \frac{\exp(-F_u \cdot F_v)}{1 - \exp(-F_u \cdot F_v)} - \sum_{v \notin N(u)} F_{v,B}, \\ & \sum_{v \in N(u)} F_{v,C} \frac{\exp(-F_u \cdot F_v)}{1 - \exp(-F_u \cdot F_v)} - \sum_{v \notin N(u)} F_{v,C}, \\ & \dots, \rangle \end{aligned}$$

Or: *for each community*, the corresponding entry to the vector  $\nabla l(F_u)$  is the one that pushes  $u$  closer to the communities in it's neighbor set  $N(u)$  and further from the communities not in its neighbor set.

# The BigCLAM Iteration

**The full update:**

$$\begin{aligned} \nabla l(F_u) = \langle & \sum_{v \in N(u)} F_{v,A} \frac{\exp(-F_u \cdot F_v)}{1 - \exp(-F_u \cdot F_v)} - \sum_{v \notin N(u)} F_{v,A}, \\ & \sum_{v \in N(u)} F_{v,B} \frac{\exp(-F_u \cdot F_v)}{1 - \exp(-F_u \cdot F_v)} - \sum_{v \notin N(u)} F_{v,B}, \\ & \dots, \rangle \end{aligned}$$

**Or Iterate:**

1. Compute gradient of  $l(F)$  with respect to (vector)  $F_u$ :  $\nabla l(F_u)$  (keeping others fixed)
2. Update the row  $F_u$  as:  $F_u^{new} = F_u^{old} + \nu \cdot \nabla l(F_u)$ . ( $\nu$  is a step size (usually small))
3. If any component  $c$  of  $F_u$  is negative ( $F_{u,c} < 0$ ), reset  $F_{u,c} = 0$ . (*Reflect*: why might this happen?)

## The BigCLAM Iteration

1. Compute gradient of  $l(F)$  with respect to (vector)  $F_u$ :  $\nabla l(F_u)$  (keeping others fixed)
2. Update the row  $F_u$  as:  $F_u^{new} = F_u^{old} + \nu \cdot \nabla l(F_u)$ .
3. If any component  $c$  of  $F_u$  is negative ( $F_{u,c} < 0$ ), reset  $F_{u,c} = 0$ .

As written, this happens to be pretty slow! We can spruce it up a little, though! The steps in vector shorthand:

$$\nabla l(F_u) = \sum_{v \in N(u)} F_v \frac{\exp(-F_u \cdot F_v)}{1 - \exp(-F_u \cdot F_v)} - \sum_{v \notin N(u)} F_v$$

**Cleanup:**  $F$  is sparse, since  $N(u)$  is usually much smaller than all nodes. This means most of the additions are in the  $\sum_{v \notin N(u)}$  sum. But we could rewrite:

$$\sum_{v \notin N(u)} F_v = \sum_v F_v - F_u - \sum_{v \in N(u)} F_v$$



# The BigCLAM

$$\begin{aligned}\nabla l(F_u) &= \sum_{v \in N(u)} F_v \frac{\exp(-F_u \cdot F_v)}{1 - \exp(-F_u \cdot F_v)} - \sum_{v \notin N(u)} F_v \\ &= \sum_{v \in N(u)} F_v \frac{\exp(-F_u \cdot F_v)}{1 - \exp(-F_u \cdot F_v)} - \left( \sum_v F_v - F_u - \sum_{v \in N(u)} F_v \right)\end{aligned}$$

# The BigCLAM

$$\begin{aligned}
 \nabla l(F_u) &= \sum_{v \in N(u)} F_v \frac{\exp(-F_u \cdot F_v)}{1 - \exp(-F_u \cdot F_v)} - \sum_{v \notin N(u)} F_v \\
 &= \sum_{v \in N(u)} F_v \frac{\exp(-F_u \cdot F_v)}{1 - \exp(-F_u \cdot F_v)} - \left( \sum_v F_v - F_u - \sum_{v \in N(u)} F_v \right) \\
 &= \sum_{v \in N(u)} F_v \left( \frac{\exp(-F_u \cdot F_v)}{1 - \exp(-F_u \cdot F_v)} + 1 \right) + F_u - \sum_v F_v
 \end{aligned}$$

# The BigCLAM

$$\begin{aligned}
 \nabla l(F_u) &= \sum_{v \in N(u)} F_v \frac{\exp(-F_u \cdot F_v)}{1 - \exp(-F_u \cdot F_v)} - \sum_{v \notin N(u)} F_v \\
 &= \sum_{v \in N(u)} F_v \frac{\exp(-F_u \cdot F_v)}{1 - \exp(-F_u \cdot F_v)} - \left( \sum_v F_v - F_u - \sum_{v \in N(u)} F_v \right) \\
 &= \sum_{v \in N(u)} F_v \left( \frac{\exp(-F_u \cdot F_v)}{1 - \exp(-F_u \cdot F_v)} + 1 \right) + F_u - \sum_v F_v \\
 &= \sum_{v \in N(u)} F_v \left( \frac{1}{1 - \exp(-F_u \cdot F_v)} \right) + F_u - \sum_v F_v
 \end{aligned}$$

# The BigCLAM

$$\begin{aligned}
 \nabla l(F_u) &= \sum_{v \in N(u)} F_v \frac{\exp(-F_u \cdot F_v)}{1 - \exp(-F_u \cdot F_v)} - \sum_{v \notin N(u)} F_v \\
 &= \sum_{v \in N(u)} F_v \frac{\exp(-F_u \cdot F_v)}{1 - \exp(-F_u \cdot F_v)} - \left( \sum_v F_v - F_u - \sum_{v \in N(u)} F_v \right) \\
 &= \sum_{v \in N(u)} F_v \left( \frac{\exp(-F_u \cdot F_v)}{1 - \exp(-F_u \cdot F_v)} + 1 \right) + F_u - \sum_v F_v \\
 &= \sum_{v \in N(u)} F_v \left( \frac{1}{1 - \exp(-F_u \cdot F_v)} \right) + F_u - \sum_v F_v
 \end{aligned}$$

What did we win?? Original RH sum:  $v \notin N(u)$  was linear in total # of nodes. Now we have just  $|N(u)|$  size updates! We can also cache/re-use the sum-over-people community scores in  $\sum_v F_v$ !

## Acknowledgments

We will implement BigCLAM in a course notebook. But there are a couple of major concerns with the algorithm that we'll touch on to open next time:

1. How do we initialize  $F$  for our gradient ascent?
2. How might we choose  $k$ ?

Next time: More on graphs: *cuts* and *partitions*, too!

Some material is adapted/adopted from Mining of Massive Data Sets, by Jure Leskovec, Anand Rajaraman, Jeff Ullman (Stanford University) <http://www.mmds.org>

